

Anh Dung CONG

0373-716-958 | dungca1512@gmail.com | LinkedIn | GitHub

OBJECTIVE

As a soon-to-be graduate with 1 year of experience working at an AI company, I am eager to apply and further develop my skills in the role of an AI Engineer at a forward-thinking company. I am passionate about participating in innovative projects and contributing to the development of cutting-edge AI solutions while continuously learning and advancing my knowledge in artificial intelligence and machine learning.

EDUCATION

VNU University of Science (HUS)
Computer and Information Science
• Good academic achievement

Hanoi, Vietnam
Mar 2019 – Jun 2025

EXPERIENCE

AI/ML Systems Architect
eUp Group

Jan 2025 – Present
Hanoi, Vietnam

AI Infrastructure & MLOps ownership

Owned end-to-end AI infrastructure across the stack: provisioning (Terraform, Ansible) → containerization (Docker, Docker Compose) → orchestration (Kubernetes/GKE, Helm) → GitOps CI/CD (GitHub Actions, GitLab CI, Jenkins, ArgoCD) → ML serving.

Deployed and operated production ML services (ASR with faster-whisper/CTranslate2, TTS, Pronunciation Scoring with wav2vec2, Qwen3 embedding) on GKE and DigitalOcean for users in VN/JP/KR.

Migrated infrastructure from Docker Swarm to GKE; managed multi-cloud across GCP, DigitalOcean, OCI, and Cloudflare (named Tunnel, no-sudo, SSE buffering workaround).

Built a bare-metal homelab Kubernetes cluster from scratch (kubeadm) with Harbor private registry, ArgoCD (App-of-Apps), and kube-prometheus-stack (Prometheus/Grafana/Alertmanager).

Cost engineering: created a multi-criteria GPU sizing framework (price/region/SLA/latency) selecting GCP L4 over A100/H100 to avoid over-provisioning; performed RCA on a 3x ASR latency gap (shared vCPU + memory-bandwidth bottleneck).

Mazii Dictionary: researched and implemented a recommendation system to personalize content on Mazii's community page.

AI Engineer
AMELA Technology

Nov 2024 – Jan 2025
Hanoi, Vietnam

Research on OCR problem for Japanese handwriting.

Research on RAG problem for chatbot.

Research on the problem of automatic cosmetic eyeliner using AI.

AI Engineer Intern
FPT Smart Cloud

Mar 2023 – Nov 2024
Hanoi, Vietnam

FPT AI Enhance for Home Credit, FE Credit, MB Bank

Participated in the development and deployment of the FPT AI Enhance project.

Handled big data processing, and trained and evaluated models for the NLP task of automatic call center scoring.

FPT AI Enhance for FPT Long Chau, FPT Shop

Built and deployed the FPT AI Enhance pipeline.

Managed big data processing, model training and evaluation, and automation of processes and data integration using n8n.

Guardrails for LLMs

Developed and tested the Guardrails project to improve the safety and control of large language models (LLMs).

Virtual Assistant Chatbot for Various Organizations

Contributed to developing and deploying virtual assistant chatbot projects for organizations such as the State Securities Commission of Vietnam, Home Credit, and FE Credit.

PROJECTS

AI Gateway

2025

Personal Project

Java, Spring Boot WebFlux, Python, FastAPI, Docker, Redis

Built a production-grade AI Gateway combining Java (Spring Boot WebFlux) gateway and Python (FastAPI) worker for unified AI request management.

Implemented multi-provider routing (OpenAI, Gemini, Claude, local LLM) with fallback, retry logic, circuit breaker, and rate limiting (Bucket4j).

Developed a Python worker supporting local LLM inference, embeddings, and a RAG pipeline with document indexing and search.

Integrated Redis caching, Prometheus monitoring, and deployed the full system via Docker Compose.

Research Agent

2025

Personal Project

Python, LangChain, LangGraph, Gemini API, MCP, A2A Protocol

Built an AI-powered research assistant for automated literature research, paper discovery, and report generation.

Designed a multi-agent architecture (Search, Paper, Synthesis agents) communicating via A2A Protocol with MCP Gateway integration.

Implemented multi-source search (Web + ArXiv), iterative research expansion, paper summarization, and structured Markdown report generation with citations.

Exposed as an MCP server for integration with Claude Desktop, Cursor, and other MCP clients.

Whisper Japanese ASR Finetune

2025

Personal Project

Python, PyTorch, HuggingFace, LoRA, Kaggle, GitHub Actions

Fine-tuned OpenAI Whisper Tiny for Japanese speech recognition using LoRA on ReazonSpeech dataset.

Built full MLOps pipeline with CI/CT/CD using GitHub Actions, Kaggle for training, and HuggingFace Hub for model hosting.

Implemented quality gate for model promotion and exported to faster-whisper (INT8) for efficient inference.

Pronunciation Scoring API

2025

eUp Group

Python, FastAPI, wav2vec2, GPU (GCP L4 / RTX 4000 Ada)

Designed a FastAPI service using wav2vec2 CTC forced alignment with a G2P caching layer and GPU micro-batching, targeting sub-3s latency for 20 concurrent users.

Sized the GPU workload (wav2vec2 315M params, FP16) to fit L4 / RTX 4000 Ada, avoiding wasteful A100/H100 provisioning.

Internal Embedding Service

2025

eUp Group

Qwen3-Embedding-4B, FastAPI, Docker Compose, RTX 4080

Deployed a self-hosted OpenAI-compatible `/v1/embeddings` service (Qwen3-Embedding-4B on RTX 4080) to fully replace the OpenAI Embedding API for internal workloads.

Built a RAG pipeline for automated code review: harvest knowledge → embed → LanceDB → inject into Qodo PR Agent, with an exit-code gate in GitLab CI that fails the pipeline on vector-store corruption.

TECHNICAL SKILLS

Programming: Python, Java, Scala, Bash, Big Data (Spark)

Cloud: GCP (~14mo, GKE/Compute/IAM), AWS (EC2/S3/Lambda/Athena), DigitalOcean, OCI, Cloudflare

DevOps/Infra: Docker, Docker Compose, Kubernetes (GKE + bare-metal kubeadm), Helm, Terraform, Ansible, ArgoCD, Cloudflare Tunnel, Harbor

CI/CD & Observability: GitHub Actions, GitLab CI, Jenkins, GitOps, Prometheus, Grafana

Libraries/Frameworks: PyTorch, TensorFlow, FastAPI, Spring Boot (WebFlux), LangChain, LangGraph, faster-whisper/CTranslate2, wav2vec2

Tools: Git, Linux, Redis, MongoDB, LanceDB, Apache Kafka, n8n, HuggingFace, Kaggle

Office: Microsoft Word, Excel, PowerPoint

English: B1, able to communicate and read technical documents

Soft Skills: Communication, Presentation, Teamwork

Research: Able to understand and implement scientific papers

CERTIFICATIONS

Mathematics for Machine Learning and Data Science Specialization (2023)

Machine Learning Specialization (2023)

HONORS & AWARDS

Published 3 Whisper-based Japanese ASR models on HuggingFace Hub ([dungca](#)); LoRA variant reached 40+ downloads.

AI Gateway project submitted to the Anthropic “Claude for Open Source” program.

Second Prize in Scientific Research Project (Faculty level) – May 2022

Third Prize in Scientific Research Project (Faculty level) – May 2024

REFERENCES

Tran Trung Hieu – Senior AI Engineer, NLP team member, FPT Smart Cloud – 0945-352-916 – hieutt125@fpt.com

Nguyen Duc Quang – Senior Developer, Team leader, AMELA Technology – quang.nguyen@amela.vn